EDGAR-SEARCH-MCP · SEC EDGAR AS A READ-ONLY MCP SERVER

The architecture in diagrams

Nine tools: companies, reported facts, filings, their text.
Every figure carries the filing it came from.

FRAMING · THE WHOLE SYSTEM · THE NINE TOOLS · HOW THEY CHAIN · ONE CALL · SEC ACCESS · STATEMENTS · FILING TEXT · SEGMENTS · BOUNDS · ERRORS · EXAMPLES · LIMITATIONS · SETUP · TESTS

VERSION 0.5.0 · OCTOBER 17, 2026 · JIBARIX · GITHUB.COM/JIBARIX/EDGAR-SEARCH-MCP · MIT LICENCE

A small server that lets a model read the SEC

An MCP client (Claude Code, Claude Desktop, any agent) starts it as a subprocess and calls its tools. It reads; it never writes anything at the SEC.

the problem

Models recall figures; filings state them

A model asked for Costco's revenue answers from memory, which is old or wrong. The answer is on file at the SEC: XBRL facts, filings, exhibits. They are free and public, but scattered over four hosts with their own formats, rate limit and identification rule.

the user

An agent, or a person through one

Any MCP client. The model picks a tool, passes plain arguments (a ticker, a form type, a URL it was given) and gets JSON back. Errors start with a fixed code, so the model can tell a typo from an outage.

what it returns

Facts with their source

Statements by period, single concepts, filing lists, a filing's documents, a document's text page by page, segment revenue, full-text search hits, candidate peers. Every statement value and concept row names its accession number, form and filing date.

rule 1 · read-only

Nothing is changed anywhere

All nine tools are marked `read_only_hint` and `open_world_hint`. Only GET requests, only to sec.gov hosts, only to URLs the server builds or checks. Redirects are not followed.

rule 2 · identify, then pace

The SEC's access rules, enforced

No request leaves without `EDGAR_USER_AGENT` (a name and email). At most 8 requests a second per process, under the SEC's 10; `Retry-After` is honoured; a 403, or a pause over 60 s, is read as rate limiting.

rule 3 · as filed, or say so

No figure is invented

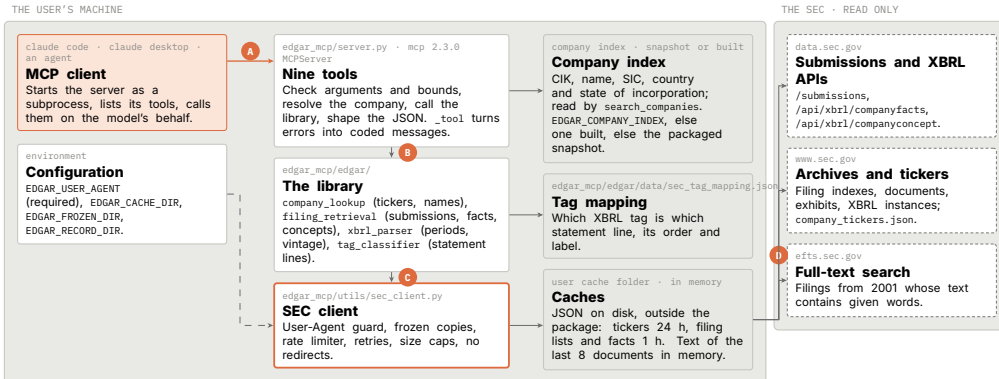
Values are the reported ones. A quarter the company did not report alone is derived from two filings and marked `derived`. When nothing fits, the result says why (an empty reason) instead of guessing.



CODE `edgar_mcp/server.py` (the nine tools), `edgar_mcp/edgar/` (lookup, filings, XBRL), `edgar_mcp/utils/sec_client.py` (every request to the SEC).
CHECKED Five test scripts, all passing on 2026-10-17: two offline, frozen replay, periods against 22 companies, the server over stdio.

The whole system on one page

One process on the user's machine, started by the MCP client. Four SEC hosts, one local index, three caches. Follow the letters.



A the client calls a tool over stdio (JSON-RPC) · **B** the tool resolves the company and calls the library · **C** every request goes through the one client
D a cache miss goes to the SEC, paced and identified; a hit never leaves the machine. `search_companies` reads only the local index.

CODE CHECKED — `edgar_mcp/__main__.py` starts `mcp.run(transport="stdio")`; `_tool` runs each tool in a worker thread; `utils.sec_client.fetch_response` is the one way out. `test_mcp_stdio.py` starts the server as a client would and calls every tool through the protocol.

The nine tools and what each is for

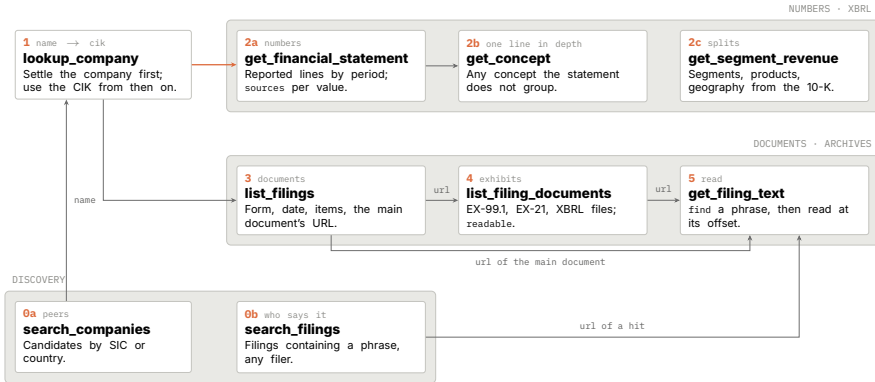
A company is a CIK, an exact ticker, an exact name or a unique name without its suffix (Apple → Apple Inc.); anything else is an AMBIGUOUS_COMPANY error that lists the candidates.

TOOL	PURPOSE	KEY ARGUMENTS	READS
`lookup_company`	Turn a name or ticker into a CIK. Says how it matched: ticker, name, shared name, prefix or fuzzy (up to 5 candidates), none.	`query`	company_tickers.json
`get_financial_statement`	One statement (BS, IS, CF, EQ, CI or ALL) by period, each value with its filing. Annual, single quarters or year to date.	`statement_type`, `period_type`, `num_periods`, `as_of`	companyfacts
`get_concept`	Every filed value of one XBRL concept, newest first: what the statement leaves out (cover share count, any company tag).	`concept`, `taxonomy`, `unit`, period range, `limit`, `offset`	companyconcept
`list_filings`	A company's filings of one form, newest first, with the main document's URL and, for 8-Ks, the item numbers.	`form`, `as_of`, `limit`, `include_amendments`	submissions (and older pages)
`list_filing_documents`	Every document of one filing: main document, exhibits, data files; which ones `get_filing_text` can read.	`url` of any document in the filing	the filing index
`get_filing_text`	A document as plain text, a page at a time, or the places a phrase occurs, with offsets to read them.	`url`, `find`, `start`, `max_chars`	the document
`get_segment_revenue`	Revenue by business segment, product or service, and geography, from the latest 10-K (else 20-F).	`as_of`	the filing's XBRL instance
`search_filings`	Filings anywhere at the SEC whose text contains given words; phrase in quotes, OR for alternatives.	`query`, `forms`, `filed_from`, `filed_to`, `offset`, `limit`	efts full-text search
`search_companies`	Candidate peers by SIC code, industry, country or name; discovery only, not evidence.	`sic`, `industry`, `country_inc`, `name_substring`	local index (no request)

CODE CHECKED One function each in `edgar_mcp/server.py`, decorated `@_tool` and registered with `ToolAnnotations(read_only_hint=True, open_world_hint=True)`. `test_mcp_stdio.py` lists the tools and calls all nine; `test_cases.py` checks the arguments of each.

How the tools hand work to each other

Three paths: numbers, documents, discovery. A URL is never typed by the model; it comes from a tool that built it.

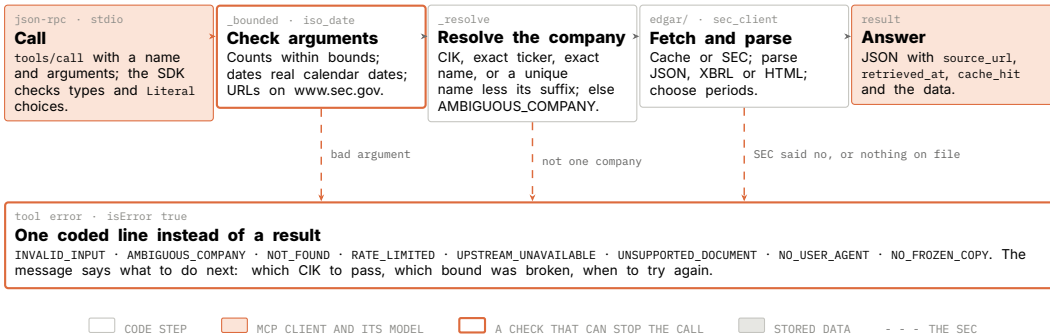


Typical questions: **"What did Costco earn?"** 1 → 2a · **"What did the earnings release say?"** 1 → 3 (8-K) → 4 (EX-99.1) → 5
"Who are the peers?" 0a, or 0b on a phrase such as "membership fee", then 1 and 2a for each candidate.

CODE CHECKED The outputs carry the inputs of the next tool: `cik` from `lookup_company`, `url` from `list_filings` and `search_filings`, `next_start` and `next_offset` for paging. `test_mcp_stdio.py` runs the chain lookup → list_filings → list_filing_documents → get_filing_text.

One call, from the model's request to its answer

The tool runs in a worker thread, so a slow SEC response does not block the server. Every failure leaves as one line that starts with its code.

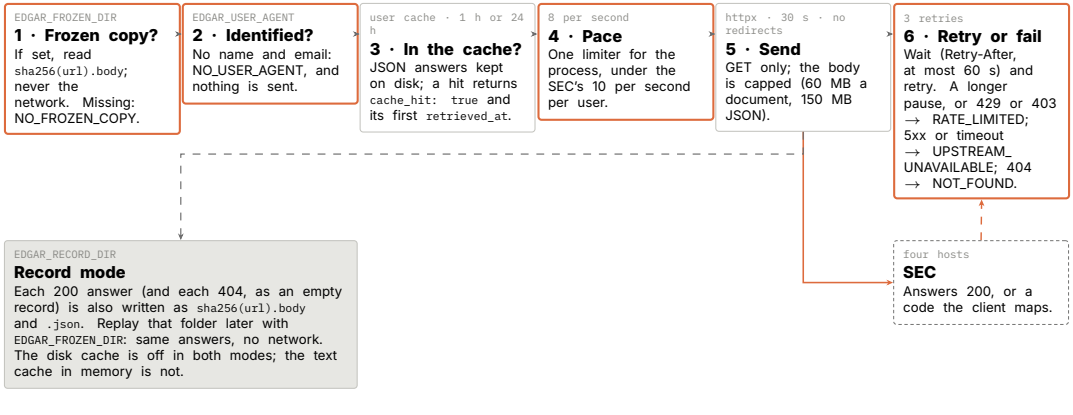


CODE `_tool` wraps each function: `asyncio.to_thread`, `SecError` → `ToolError("CODE: message")`, `ValueError` → `INVALID_INPUT`.

CHECKED `test_cases.py` checks each error path offline; `test_mcp_stdio.py` checks the codes arrive through the protocol, including NO_USER_AGENT.

Every request to the SEC passes the same six steps

Nothing in the server opens a connection of its own. The order matters: a frozen copy is checked before the User-Agent, so replays work offline and anonymously.

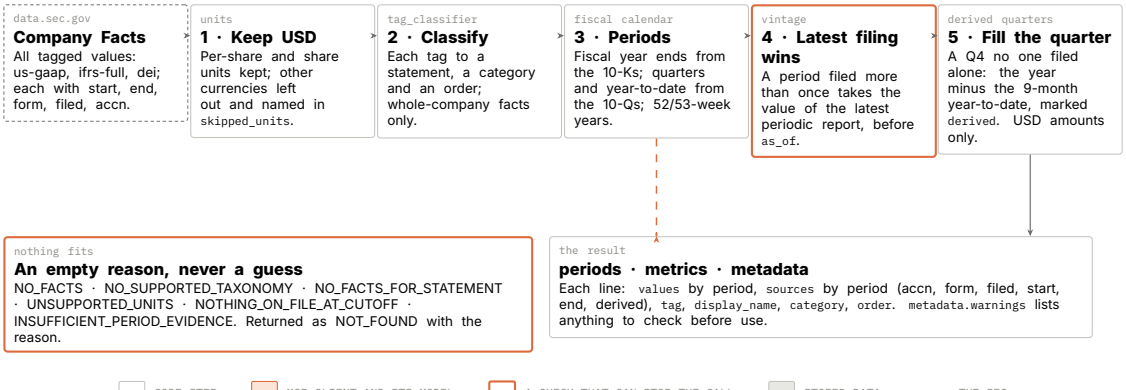


CODE utils/sec_client.fetch_response and get_json; config/settings.py (8 per second, 3 retries, 30 s timeout); utils/cache.py.

CHECKED test_cases.py: pacing, Retry-After, 403, 5xx, size cap, no User-Agent; test_tools_offline.py: a long pause on a 503, malformed frozen copies; test_frozen.py: record then replay with no network.

From Company Facts to a statement by period

The SEC's Company Facts feed holds every value a company ever tagged, often filed three or four times. Five steps pick one value per line and period.



CODE edgar/xbrl_parser.parse_company_facts; edgar/tag_classifier.classify_with_report with sec_tag_mapping.json; empty reasons in server._EMPTY_REASONS.
CHECKED test_periods.py: 22 companies, 243 annual, 46 quarterly and 46 year-to-date figures against the raw 10-K facts; 0 problems on 2026-10-17.

Reading a filing document: checked address, plain text, pages

A 10-K runs to a million characters. The tool returns a page (12,000 characters by default) or the offsets where a phrase occurs, and the SHA-256 of the file it read.



list_filing_documents

The filing index, row by row

Takes any document URL of the filing, reads the index page and lists every document: sequence, type (8-K, EX-99.1, GRAPHIC, XML), size, URL, and `readable` (step 3 would accept it).

provenance

What a reader can cite

`url`, `sha256` of the file as sent, `extraction` (the rules' version), `retrieved_at`, `total_chars`, `start`. The same text, by the same rules, is open to programs through `document_text`.

CODE STEP · MCP CLIENT AND ITS MODEL · A CHECK THAT CAN STOP THE CALL · STORED DATA · - - - THE SEC

CODE `server._document_url`, `document_text` (extraction `tags-stripped-v2`), `_document_text` (in-memory cache), `get_filing_text`, `list_filing_documents`.

CHECKED `test_cases.py`: URL rejections, content types, extraction, `find` paging; `test_mcp_stdio.py`: reads a real 10-K and an exhibit.

Segment revenue, full-text search and candidate peers

Segment splits are not in Company Facts: the tool reads the 10-K's own XBRL instance. Search and peers are for finding candidates, which the other tools then confirm.

get_segment_revenue

Revenue by axis, from the filing itself

1. `list_filings` finds the latest 10-K (else 20-F) before `as_of`. 2. Its index names the XBRL instance. 3. Annual figures of standard-taxonomy concepts on three axes are kept, in one currency (the total's) and one period; the rest go to `skipped` with the reason. 4. Each member with its value, concept and share of the members listed; `concepts` names the concepts read per axis.

caveat

Members can overlap

A company may tag a total member (`ProductMember`) next to its parts; segment sales may include sales between segments. `sum_of_members` shows it; shares are of the members listed, not of revenue.

search_filings

Who wrote these words, and where

The SEC's full-text search over filings since 2001: all words required, a phrase in double quotes, OR between alternatives; filtered by form and filing dates. Each hit gives the filer, its SIC, the document URL and whether it is `readable`. 100 hits a page; the SEC returns at most 10,000.

caveat

Words, not meaning

`meaning: "the words are in the document; proximity and meaning are not checked".` A hit is a place to read, not a finding.

search_companies

Candidates from a local index

No request leaves the machine. Filters are combined with AND: exact 4-digit SIC, industry, country or state of incorporation, part of the name. The index is the packaged snapshot (periods to 2025-10-31, build date not recorded) or one built from the SEC's Financial Statement Data Sets; `index_origin` says which. `use: "candidate discovery only".`

caveat

Old and partial

Companies that never filed financial statement data, or changed SIC since the index was built, are missing or stale. Confirm each with `lookup_company`.

CODE STEP · MCP CLIENT AND ITS MODEL · A CHECK THAT CAN STOP THE CALL · STORED DATA · - - - THE SEC

CODE `get_segment_revenue` (XBRL instance, `_SEGMENT_AXES`, `_REVENUE_CONCEPTS`), `search_filings` (efts, 100 a page), `search_companies` (`company_classifier.load_index_with_meta`).

CHECKED `test_tools_offline.py` (currencies, units, periods, namespaces, malformed search answers, the 10,000 ceiling, index precedence) and `test_mcp_stdio.py` (each tool on a live company).

Parameters, defaults and bounds

Dates are YYYY-MM-DD and must exist. `as_of` means “as the record stood that day”: anything filed later is ignored.

TOOL	PARAMETER	DEFAULT · BOUND	NOTE

`get_financial_statement`	`statement_type`
	`period_type`
	
	`num_periods`
		BS IS CF EQ CI ALL
		annual · quarterly · ytd
		
		4 · 1–40
			ALL returns every statement in one call.
			ytd: fiscal year to date at each quarter end; with annual, gives a trailing twelve months.
			Periods, newest first.

`get_concept`	`taxonomy`
	`limit, offset`
	`unit, period_end_from/to`
		us-gaap · ifrs-full · dei
		100 · 1–500
		none
			dei holds cover-page items such as the share count.
			`next_offset` pages through `total_rows`.
			`units_available` lists the choices.

`list_filings`	`form, limit`
	
	`include_amendments`
		10-K · 10 · 1–100
		
		false
			Older submission pages are read when needed; `complete` says if all were.
			10-K/A and the like.

`get_filing_text`	`max_chars` `find`	12,000 · up to 40,000 up to 200 characters	A page; `next_start` continues. Case-insensitive; 40 matches a call, then `next_start`.
`search_filings`	`limit, offset` `forms`	25 · 1–100 · offset to 9,999 10-K	The SEC returns at most 10,000 hits; `next_offset` is null at the end. One or several, comma-separated; amendments come with their form.
`search_companies`	`limit`	25 · 1–200	`truncated` says more matched.
all company tools	`cik_or_ticker`	—	CIK (any zero padding), exact ticker, or exact name.

CODE `_bounded(name, value, low, high)` and `iso_date` at the top of each tool; out of range is INVALID_INPUT, never silently clipped.
CHECKED `test_cases.py` tries the edge of each bound.

Error codes, and what to do about each

A tool error is one line, `CODE: message`. A result with nothing in it is not an error: an empty list, or a reason in `metadata.empty_reason`.

CODE	MEANS	WHAT TO DO
`INVALID_INPUT`	An argument out of bounds, a date that does not exist, a URL that is not an EDGAR document.	Fix the argument; the message names it and its bound.
`AMBIGUOUS_COMPANY`	The company text is not an exact ticker, name or CIK.	Call again with one of the CIKs listed in the message.
`NOT_FOUND`	The SEC has nothing at that address, or no figures fit (with the empty reason).	Check the concept, form or period; read the reason.
`RATE_LIMITED`	The SEC answered 429 or 403 after three retries, or asked for a pause over 60 s (a 503 with `Retry-After`).	Wait; the message says how long is left.
`UPSTREAM_UNAVAILABLE`	5xx, a timeout or a broken connection after three retries, or an answer of the wrong shape.	Try again later.
`UNSUPPORTED_DOCUMENT`	A PDF, an image, a file over 60 MB, a 10-K without inline XBRL, or revenue in two currencies with no total.	Pick another document; for segments, read the text.
`NO_USER_AGENT`	`EDGAR_USER_AGENT` is not set; nothing was sent.	Set it to a name and email, then restart the server.
`NO_FROZEN_COPY`	Replay mode, and this URL was never recorded.	Record it with `EDGAR_RECORD_DIR`, or leave replay mode.
`INDEX_UNAVAILABLE`	The company index named or built cannot be read.	Fix or remove `EDGAR_COMPANY_INDEX`, or rebuild.

CODE `SecError` subclasses in `utils/sec_client.py`; `_tool` in `server.py` prefixes each message with its code.
CHECKED `test_cases.py` and `test_tools_offline.py` (every code, RATE_LIMITED offline), `test_mcp_stdio.py` (codes over the protocol), `test_frozen.py` (NO_FROZEN_COPY).

Example 1: from a name to an income statement

“What did Costco earn last year?” Two calls. Every value says which filing it came from.

```
lookup_company(query="Costco")
match_type: "prefix"
matches: [{cik: "0000909832", ticker: "COST",
           name: "COSTCO WHOLESALE CORP /NEW"}]
```

“Prefix” means a candidate, not a certainty: the model checks the name, then uses the CIK or ticker.

```
get_financial_statement(cik_or_ticker="COST",
  statement_type="IS", period_type="annual", num_periods=2)
periods: ["2026-08-30", "2025-08-31"]
source_url: .../companyfacts/CIK0000909832.json
metrics.<line>.sources["2026-08-30"]:
  {accn: "0000909832-26-000093", form: "10-K",
   filed: "2026-10-07", start: "2025-09-01", derived:
false}
```

the lines, as returned (millions of USD)

Costco, fiscal years to August

LINE	FY2026	FY2025
Net revenue	303,154	275,235
Cost of revenue	264,279	—
SG&A	27,190	—
Operating income	11,685	10,383
Income before income taxes	12,251	—
Income tax	3,025	—
Net income	9,226	8,099
Diluted EPS (USD)	20.76	18.21

The tag of each line is in `tag` (Net revenue is `us-gaap:RevenueFromContractWithCustomerExcludingAssessedTax`). Values come back in full units: `303154000000`. Line names here are plain English; — marks values not shown.

CODE Real calls on 2026-10-17, shortened: values in USD as returned, here written in millions.
CHECKED The same figures appear in Costco’s 10-K for fiscal 2026 (accession 0000909832-26-000093).

Example 2: a derived quarter, and a figure no statement holds

Costco never files its fourth quarter on its own. The tool derives it from two filings and says so. The share count lives on the cover page, under dei.

```
get_financial_statement("COST", "IS", period_type="quarterly")
Net revenue, by quarter end (millions):
   2026-08-30       95,723    derived
   2026-05-10       70,527
   2026-02-15       69,597
   2025-11-23       67,307    sum 303,154 = the 10-K year
sources["2026-08-30"]:  {derived:  true,
  cumulative:  303,154 (10-K 0000909832-26-000093, year),
  minus:  207,431 (10-Q 0000909832-26-000051, 36 weeks)}
```

The quarter is the year less the year to date at the third quarter. Both filings are named, so a reader can check the subtraction.

```
get_concept("COST", "EntityCommonStockSharesOutstanding",
  taxonomy="dei", limit=2)
label:  "Entity Common Stock, Shares Outstanding"
unit:  "shares"   total_rows:  68   next_offset:  2
facts:  [
 {end:  "2026-09-29", val:  443129349,
  form:  "10-K", filed:  "2026-10-07", fp:  "FY"},
 {end:  "2026-05-27", val:  443478804,
  form:  "10-Q", filed:  "2026-06-03", fp:  "Q3"}]
```

Newest first. `next_offset` reads the next two; `as_of` and the period range narrow the rows; `scope: "whole-entity values only"`.

CODE Real calls on 2026-10-17, shortened.
CHECKED `test_periods.py` checks this rule on 22 companies: the four quarters sum to the year within 0.05%.

Example 3: what did the earnings release say?

The numbers of an 8-K earnings release are in its exhibit 99.1, not in the 8-K. Three calls find it and read the passage.

1 `list_filings`

The 8-Ks, newest first

```
list_filings("COST", form="8-K", limit=10)
```

```
2026-10-16      items 5.02, 9.01
2026-10-14      items 8.01, 9.01
...
2026-09-24      items 2.02, 9.01
```

Item 2.02 is "Results of operations": the earnings release. Each row carries the main document's `url`.

2 `list_filing_documents`

What the filing holds

```
list_filing_documents(url=<the 2026-09-24 8-K>)
```

```
count: 18
type          document                 readable
8-K           main document            yes
EX-99.1       costex9918-k92426.htm    yes
EX-99.2       .htm                     yes
GRAPHIC       .jpg                     no
EX-101.SCH    .xsd                     no
XML           XBRL instance            yes
```

3 `get_filing_text`

Find, then read

```
get_filing_text(url=<EX-99.1>, find="net income")
```

```
total_chars: 8708  match_count: 6
extraction: "tags-stripped-v2"
sha256: "651f6ec2..."
```

"Net income for the fourth quarter was $2.998 billion, $6.75 per diluted share, compared to $2.610 billion..."
"Net income for the fiscal year was $9.226 billion, $20.76 per diluted share..."

Each match has an `offset`; `start=<offset>` reads the paragraph around it.

The release's $9.226 billion matches the 10-K's 9,226 million in Example 1: two sources, one figure.

CODE Real calls on 2026-10-17, shortened.
CHECKED `test_mcp_stdio.py` runs the same chain on a live company each time it runs.

Example 4: where the revenue comes from, and who else says “membership fee”

Segments from the 10-K’s XBRL instance; candidates from full-text search and from the local index. The last two find names to check, not answers.

```
get_segment_revenue("COST")
form: "10-K", filing_date: "2026-10-07"
fiscal_year: 2025-09-01 to 2026-08-30
source_url: .../cost-20260830_htm.xml
unit: "USD"  skipped: []  frozen: false
total_revenue: 303,154 (millions)
business_segments:
    United States              219,823    0.7251
    Other International         42,770    0.1411
    Canada                      40,561    0.1338
products_and_services:
    Product                    297,247    0.4951 (a total)
    Foodand Sundries           116,788    0.1945
    Non Foods                   77,375    0.1289
    Other, Fresh Food, ...
```

```
search_filings('"membership fee"',
  forms="10-K", filed_from="2025-01-01",
  filed_to="2025-12-31", limit=2)
total_hits: 46  total_is_exact: true
next_offset: 2
Exceed World, Inc.      SIC 8200
BJ's Wholesale Club     SIC 5331
  bj-20250201.htm  readable
```

Hits include filers far from retail: the phrase alone does not make a peer.

```
search_companies(sic="5331", limit=3)
index_origin: "snapshot in the
package"
index_built: "unknown"
use: "candidate discovery only"
COSTCO          US · WA
PRICESMART      US · DE
MINISO          US · KY
truncated: true
```

SIC 5331 is “Variety Stores”. BJ’s, found by search, is in it too.

CODE Real calls on 2026-10-17, shortened.

CHECKED The segment total equals the 10-K’s revenue: 219,823 + 42,770 + 40,561 = 303,154 million.

Example 5: four errors, each telling the model what to do

The message is written for the model that made the call: it names the bad argument, or the CIKs to choose from, or the setting to fix.

```
get_financial_statement("Costco", "IS", as_of="2024-13-01")
INVALID_INPUT: '2024-13-01' is not a calendar date
```

```
list_filings("Apple Hosp")
AMBIGUOUS_COMPANY: 'Apple Hosp' is not an exact ticker or
company name.  Call again with the CIK of one of:  PPL Corp
(PPL, CIK 0000922224); Apple Inc.  (AAPL, CIK 0000320193)
```

```
get_concept("COST", "NoSuchConcept")
NOT_FOUND: nothing at https://data.sec.gov/api/xbrl/
companyconcept/CIK0000909832/us-gaap/NoSuchConcept.json
```

```
list_filings("COST")    # server started without EDGAR_USER_AGENT
NO_USER_AGENT: set EDGAR_USER_AGENT to your name and email,
e.g.  "Jane Doe jane@example.com"; the SEC asks every client
to identify itself (<the URL it would have read>)
```

Note the second case: "Apple Hosp" (for Apple Hospitality REIT, ticker APLE) gets two fuzzy candidates, neither the one meant. The error does not pick one; the model must. `lookup_company("APLE")` settles it. On startup without a User-Agent the server also prints a warning to stderr, which most MCP clients show in their log.

CODE Real calls on 2026-10-17; the messages are verbatim, but for the URL at the end of the last one.
CHECKED `test_cases.py` and `test_mcp_stdio.py` assert each code.

Limitations: what the server does not do

Each is a boundary chosen on purpose, and each is written into the tool's description so the model sees it before calling.

statements

Reported facts, not the statement as printed

Whole-company facts only: lines split by share class, segment or geography, and company-specific (extension) tags are left out; read them in the text. USD only; other currencies are named in `skipped_units`, so most 20-F and 40-F filers in euros or yen return little.

classification

A mapping, not the company's layout

Lines are placed by a fixed tag mapping, not by the filing's presentation linkbase. Order and sections are approximate, and tags are matched by local name; reviewed rules correct the known errors (next page). `metadata.unclassified` counts what was left out.

text

Plain text, and only text

Tables become tab-separated lines; layout and footnote marks are lost. PDFs, images and files over 60 MB are refused. No OCR. Full-text search starts in 2001; filings from before EDGAR are not there.

names

Today's tickers and names

`company_tickers.json` lists current tickers. A company renamed, merged or delisted may not resolve by name; use its CIK. Fuzzy matches can be wrong (Example 5).

freshness

An hour behind, at most

Filing lists and facts are cached for an hour and tickers for a day; `cache_hit` says when an answer is from the cache. The company index behind `search_companies` is the packaged snapshot until one is built.

scale and scope

One user, one machine, read only

8 requests a second per process; two servers on one machine share the SEC's 10 per second between them. No bulk download, no writing, no stock prices, no insider (Forms 3/4/5) or 13F parsing beyond reading the text.

Not yet built: an output schema. Every tool returns structured JSON, but its shape is described in words (tool descriptions, README), not in a machine-checkable schema. A program reading the answers directly should check the fields it uses; names may change between versions.

CODE The tool docstrings and README.md (Limitations) state each of these to the model and the user.
CHECKED Not tested as failures: these are the edges of what the server claims.

The Costco quirks: what was fixed, what is kept on purpose

The headline figures were always right. Four mapping errors are now corrected by reviewed rules; three behaviours stay, because changing them would hide what the company tagged.

WHAT CAME BACK IN 0.4.0	NOW	WHY
`us-gaap:InterestExpense` (145 million) in category “Revenue”.	Category “Income”, label “Interest expense”.	An expense; order 8 keeps a bank’s interest lines together.
`IncomeLossAttributableToParent` labelled “Income before income taxes”.	“Income (loss) attributable to parent, before tax”.	The taxonomy’s own meaning.
`InterestAndOtherIncome` labelled “Interest income”.	“Interest and other income”.	The tag holds more than interest.
Dividends, restricted stock issued and unrealized debt gains in the income statement.	Dividends and restricted stock in EQ; the debt gain left out, its reason in `metadata.unclassified`.	Equity movements; the gain can be net income or OCI. Lease cost stays: it is a reported expense.
`us-gaap:Revenues` null; `RevenueFromContract…` holds 303,154.	*Kept:* two lines, nulls shown.	Two tags can differ in scope; the model picks by `tag`.
`ProductMember` (297,247) listed beside its parts.	*Kept:* listed as tagged; `note` warns.	A total member is the company’s choice; guessing which is a total could drop a real member.
Label “Foodand Sundries”.	*Kept:* built from the member name.	The `member` field is exact; labels from the filing’s label linkbase are a possible next step.

CODE CHECKED Seen in the Costco calls of 2026-10-17 (Examples 1 and 4), before and after the fixes of version 0.5.0; causes in `sec_tag_mapping.json` and in Costco’s own member names. `test_tools_offline.py` (classifier and segment cases) holds each row’s current behaviour.

Install, configure, connect

Python 3.11 or later. Three direct dependencies, pinned: httpx 0.28.1, jsonschema 4.26.0, mcp 2.3.0; `requirements.lock` holds the hashed closure.

1 install

An ordinary install in a virtual environment

```
git clone https://github.com/jibarix/edgar-search-mcp
cd edgar-search-mcp
python -m venv .venv
.venv/Scripts/python -m pip install .     # or bin/ on macOS, Linux
```

The package carries the tag mapping and the index snapshot; the cache and any index you build go to your user folders, never next to the code.

3 connect

Claude Code

```
claude mcp add edgar-search \
  -e EDGAR_USER_AGENT="Jane Doe jane@example.com" \
  – <repo>/.venv/Scripts/python -X utf8 -m edgar_mcp
```

Claude Desktop: the same command, arguments and `env` in `claude_desktop_config.json` (README).

2 configure

Environment variables

`EDGAR_USER_AGENT`	**Required.** A name and email the SEC can reach. No default: without it nothing is sent.
`EDGAR_CACHE_DIR`	Where the JSON cache lives; default the user cache folder (`%LOCALAPPDATA%`, `~/Library/Caches`, XDG).
`EDGAR_COMPANY_INDEX`	A company index to use instead of the built or packaged one.
`EDGAR_RECORD_DIR`	Record each answer (200, or 404 as empty) as a frozen copy here.
`EDGAR_FROZEN_DIR`	Replay frozen copies from here; never touch the network.

Record and replay make a run repeatable: the same calls give the same answers next month, for a test or an audit.

CODE `edgar_mcp/config/constants.py` (`USER_AGENT`), `edgar_mcp/config/settings.py` (`_user_dir`), `edgar_mcp/__main__.py`, `pyproject.toml` (script `edgar-search-mcp`).
CHECKED `test_mcp_stdio.py` starts the server the way the command below does.

Five tests, and where the project stands

Offline checks run anywhere in seconds. The live checks read the SEC, pacing themselves, and keep what they fetched for the next run.

TEST	NETWORK	WHAT IT CHECKS	2026-10-17
`test_cases.py`	none	Argument bounds, company resolution, URL checks, text extraction and `find`, segment axes, search paging, the client's pacing, retries, Retry-After, size caps, and that no request is sent without a User-Agent.	45 checks, 0 problems
`test_tools_offline.py`	none, guarded	The tools against made-up SEC answers: malformed search answers and the 10,000 ceiling, retries and long pauses, frozen copies, history errors, segment currencies, units and periods, classification rules, index precedence and builder, user folders.	86 checks, 0 problems
`test_frozen.py`	records once	Records a set of calls with `EDGAR_RECORD_DIR`, replays them with `EDGAR_FROZEN_DIR` and the network blocked, compares; a missing copy gives NO_FROZEN_COPY.	0 problems
`test_periods.py`	cached	22 companies with different fiscal calendars: annual periods and values against the 10-K rows of the raw Company Facts (not the 10-K documents), four quarters summing to the year, year to date at the year end.	243 + 46 + 46 figures, 0 problems
`test_mcp_stdio.py`	live	Starts the server as a client does, lists the nine tools, calls each, follows the document chain, checks error codes over the protocol, and that a server with no User-Agent answers NO_USER_AGENT.	0 problems

Status. Version 0.5.0, MIT licence; the fixes are not yet committed or published. The final review's six findings (segment units, search continuation, replay timestamps, index date, empty search responses, historical-file errors) are fixed, with the classification corrections, a command to build the company index and per-user folders; each has an offline test. **Possible next step:** segment member labels from the filing's label linkbase.

CODE `tests/`: run each from the repository root, `python -X utf8 tests/<name>.py`; the live ones need `EDGAR_USER_AGENT`.
CHECKED All five, 0 problems, on 2026-10-17 (Windows 11, CPython 3.13).